

Your Viewing Guide: Weekend Catch-Up

Profile: daniel-agentic-coding-skills-founder · Jun 22-25, 2026 · 6 new videos

You asked: "catch up on everything I missed since the weekend - Raj on Fable's return, Nate is on a roll - then tell me what to watch tonight." Curated through your lens (agentic-coding craft, Claude skills, models, the founder/consulting path); a different reader would rank these differently. Every orange timestamp is a click-to-jump deep-link from the actual mindmap.

If you only have 30 minutes, watch these 2

The Doing Got Cheap. Now What? | Claude Fable 5 Changes Work

natebjones · 2026-06-23

Your top watch. The job-based **"gives"** and **Model Manager** reframe is the strategic spine of your dark-factory / autonomous-run work *and* a ready-made consulting pitch: what to hand a frontier model whole, versus what stays human judgment. Pairs with your "judgment is the moat" thread.

6:45 Asks vs Gives · **12:52** the Model Manager role · **17:02** whole-job specs + review trails

Anthropic Will Bring Back Fable 5 Differently

ramjad · 2026-06-25

The Fable news you flagged - with an operating lesson. Beyond the weekly-limits + /usage-credits mechanics you'll feel directly, Raj's real point is **targeting**: rank with cheap models, spend premium tokens only on high-impact work. That is exactly your /loop cost discipline, externally validated.

0:35 weekly limits + credit system · **3:06** impact/opportunity targeting · **10:04** human ranking as the moat

Pillar 1 - Agentic-Coding Craft & Skills

I Stopped Prompting AI One Task At A Time

natebjones · 2026-06-24

Loops and **loops-of-loops** as the unit of work - the conceptual backbone of your /loop and dark-factory harnesses. The signals -> shared-context -> control-layer model maps cleanly onto how you already orchestrate agents; the **attention-budgeting** framing is a clean way to describe quiet autonomous runs.

1:30 prompt vs loop vs loop-of-loops · **5:53** attention budgeting: notify on change · **10:57** safe-action framework

Pillar 2 - AI Strategy & Models

Qwen-AgentWorld: The World Model for RL Environments

samwitteveenai · 2026-06-25

Keeps your model radar current and feeds the **eval** thread: an open-weights world model (397B, 35B MoE released) that **replaces slow physical sandboxes** for RL trajectory generation. The LLM-judge + rule-verifier anti-reward-hacking setup rhymes with your grounded-eval instincts.

0:05 91.2 on AgentWorldBench · **4:25** replaces physical sandboxes for RL · **10:20** LLM-judge + rule verifiers

Pillar 3 - Startup / Founder / Consulting

Claude Code + AWS Bedrock = Enterprise AI

mark_kashef · 2026-06-25

A regulated-enterprise **"mission control"** you could build or sell as a consulting deliverable: in-account inference, SOC 2/HIPAA scoring, kill switches, multi-tenant isolation. **Exactly the enterprise-AI angle worth tracking for client work** - watch how the compliance scoring is wired.

0:53 the regulated-teams thesis · **12:01** Bedrock in-account inference · **20:48** 42 kill switches + audit logs

Lower signal - skim or skip

Google Lost \$2.7B In Talent This Week (natebjones, 06-22) - off your usual lane: talent migration + a Midjourney whole-body-ultrasound tangent. The only durable signal is the move map (Shazeer to OpenAI, Jumper/AlphaFold to Anthropic). Skim the first 2 minutes or skip.

Not surfaced: 14 of 19 channels were unscanned this run (YouTube quota); the post-3am sweep will fold in everinc, claude, colemedin, twist and the rest. This is the weekend slice only.